

DCEE-CausalVerse: Event-Grounded Causal Planning for Ending-Controllable Visual Storytelling

Anonymous Submission

Anonymous affiliation

Abstract

Recent text-to-image diffusion models produce high-quality single images, but visual storytelling remains difficult because image sequences must preserve identity, world state, event continuity, and emotional progression. We study ending-controllable causal visual storytelling, where a generated story must express a target ending emotion and visually justify why that emotion emerges. We propose DCEE-CausalVerse, which decomposes a target ending emotion into Desire, Conflict, Event Chain, and Ending Emotion. A DCEE-Tree planner explores multiple Desire-Conflict routes and selects the event chain that best supports the target emotion. The selected plan is canonicalized into drawable events and evidence, maintained through causal sink memory and a DCEE anchor bank, and translated into factorized diffusion controls by a decoder-encoder-decoder Butterfly Visual Controller. The controller generates Character, World, Emotion, Event, and Evidence control branches that are injected as cross-attention tokens for SDXL. Multi-candidate generation and event-evidence-aware reranking improve visual grounding and continuity. Automatic and human evaluations show stronger cross-frame consistency and clear preference for caption reflection, object/action visibility, emotion delivery, and emotion-cause explainability.

1. Introduction

Text-to-image diffusion models have made rapid progress in producing visually compelling images from natural language prompts (Ho, Jain, and Abbeel 2020). However, visual storytelling is more demanding than single-image generation because each frame must simultaneously satisfy local text-image alignment and global story consistency. The protagonist should remain recognizable, important objects should persist or reappear when causally needed, the world should evolve coherently, and the emotional tone should change in a narrative rather than arbitrary manner (Rombach et al. 2022).

Most existing visual storytelling systems treat the input as a sequence of captions and attempt to visualize each caption while maintaining history consistency. This is useful when the story is already fixed, but it does not directly

answer a more controllable question: given an initial premise and a target ending emotion, what sequence of events should happen so that the final emotion becomes believable, and how should those events be visually grounded? A direct prompt such as 'generate a sad ending' may produce a gloomy portrait, but it often fails to show the event and visual evidence that explain the sadness.

We argue that ending-controllable visual storytelling requires causal narrative planning. The final emotion should not be a surface-level style label; it should be the result of a protagonist's desire, a central conflict, and a chain of visually drawable events. Moreover, every emotional transition should have visible evidence. For example, a sad frame should not simply contain a sad face; it should show the lost object, empty hands, rainy riverbank, or broken clue that makes the sadness understandable.

To address this problem, we propose DCEE-CausalVerse. The core representation is DCEE: Desire - Conflict - Event Chain - Ending Emotion. We further extend this representation with DCEE-Tree Planning, which samples multiple Desire-Conflict trajectories before selecting the event chain that most strongly supports the target ending emotion. The selected plan is canonicalized into explicit entities, drawable events, visual evidence, and world-state cues. During generation, causal sink memory selects not just visually similar history but emotionally causal memory, while the DCEE anchor bank maintains recurring characters, evidence objects, and world anchors.

A key component of our method is the DCEE Butterfly Visual Controller. Unlike a simple prompt generator, this module follows a decoder-encoder-decoder structure. It first expands a DCEE frame into diverse visual hypotheses, then encodes them with causal memory and anchor constraints, and finally decodes the fused state into factorized Character, World, Emotion, Event, and Evidence controls. These controls are converted into cross-attention tokens for SDXL (Podell et al. 2024), enabling structured conditioning while avoiding overly long prompts and control entanglement.

In summary, our main contributions are fivefold:

1. We formulate ending-controllable causal visual storytelling, where a generated visual story must justify a target ending emotion through a causally ordered event chain.
2. We introduce DCEE-Tree Planning, which explores multiple Desire-Conflict routes and selects the event chain with the strongest causal support for the target ending emotion.
3. We propose causal sink memory and a DCEE anchor bank to preserve identity, world state, conflict state, event evidence, and emotion-cause information across frames.
4. We design a decoder-encoder-decoder DCEE Butterfly Visual Controller that maps symbolic DCEE frames into factorized Character, World, Emotion, Event, and Evidence controls.
5. We introduce event-evidence-aware reranking and retry generation using event grounding, evidence visibility, emotion visibility, and event-emotion causal consistency.

2. Related Work

2.1. Story Visualization and Subject Consistency

Story visualization methods generate one image per sentence while preserving sequence coherence (Huang et al. 2016; Li and Lukasiewicz 2022; Ahn et al. 2023). GAN- and transformer-based methods such as StoryGAN (Li et al. 2019) and StoryDALL-E (Maharana, Hannan, and Bansal 2022) initiated this task, and diffusion-based approaches such as AR-LDM (Pan et al. 2024), Make-A-Story (Rahman et al. 2023), ACM-VSG (Feng et al. 2023), StoryGen (Liu et al. 2024), Causal-Story (Song et al. 2024), and ViSTA (Dong et al. 2026) further improve auto-regressive generation through visual memory, adaptive context, causal attention, or history adapters. These methods are strong at rendering a given caption sequence, but they usually assume that the event chain is already specified rather than planning why a target ending emotion should occur.

A related line studies subject-consistent generation. IP-Adapter (Ye et al. 2023), ConsiStory (Tewel et al. 2024), StoryDiffusion (Zhou et al. 2024), One-Prompt-One-Story (Liu et al. 2025), and DreamStory (He et al. 2025) preserve recurring characters or styles across frames by sharing attention, reusing subject features, or constructing multi-subject prompts. Such identity preservation is necessary for visual storytelling, but it is not sufficient for ending-controllable storytelling: a stable protagonist may still fail to perform the intended event, interact with the required evidence object, or express an emotion for a visible reason.

2.2. Planning, LLM Guidance, and Faithfulness

LLM-guided visual story systems use language models as story directors, prompt generators, or reference resolvers. StoryGPT-V (Shen and Elhoseiny 2025), DreamStory (He et al. 2025), and SEED-Story (Yang et al. 2025) show that LLMs can organize story context and improve multi-modal story generation. In parallel, narrative planning work such as hierarchical story generation (Fan, Lewis, and Dauphin 2018), Plan-and-Write (Yao et al. 2019), PlotMachines (Rashkin et al. 2020), and reward-shaped plot generation (Tambwekar et al. 2019) shows that coherent stories benefit from explicit intermediate plans. DCEE-CausalVerse combines these ideas by using an LLM as a structured causal planner rather than a prompt writer: the target ending emotion is decomposed into Desire, Conflict, Event Chain, and Ending Emotion, then canonicalized into drawable events and visible evidence.

Faithfulness evaluation is also important because visually appealing frames can omit the intended action or causal clue (Ghosh, Hajishirzi, and Schmidt 2023; Lin et al. 2024). TIFA (Hu et al. 2023) evaluates text-to-image faithfulness with VQA-style questions, and ViSTA (Dong et al. 2026) adopts related criteria for story visualization. Our evaluation extends this direction to causal emotional storytelling by checking not only text-image alignment and identity consistency, but also whether the generated sequence makes the final emotion explainable from visible events and evidence.

3. Method

Figure 1 summarizes the overall DCEE-CausalVerse pipeline. Given an input image or text prompt, a protagonist description, and a target ending emotion, the framework builds a Global Story Bible, explores multiple DCEE routes, canonicalizes the selected route into drawable events and evidence, retrieves causal memory and anchors, and generates each frame through factorized cross-attention control. Table B1 in the appendix provides the full step-by-step pipeline summary.

3.1. Problem Formulation

Let x denote the input premise, a the protagonist description, e^* the target ending emotion, and K the number of frames. The goal is to generate an image sequence $Y_{1:K}$ and storyboard $S_{1:K}$ such that the final emotion aligns with e^* and every frame grounds the event and visible evidence that cause the protagonist's emotion:

$$(Y_{1:K}, S_{1:K}) = \mathcal{G}(x, a, e^*, K)$$

DCEE-CausalVerse: Event-Grounded Causal Planning for Ending-Controllable Visual Storytelling

From Causal Planning to Emotion-Evidential Image Generation




Figure 1. Overall pipeline of DCEE-CausalVerse

Unlike standard story visualization, the event chain is planned rather than assumed to be fully given. We represent the planned causal structure as

$$\mathcal{P}_{\text{DCEE}} = (d, c, \mathcal{E}_{1:K}, e^*),$$

where d , c , and $\mathcal{E}_{1:K}$ denote the Desire, Conflict, and Event chain, respectively.

3.2. Global Story Bible

The Global Story Bible stores protagonist identity, world setting, key objects, visual style, central desire, possible conflicts, symbolic motifs, and target ending emotion. If an initial image is provided, an image-text understanding module extracts characters, scene, objects, mood, weather, and plot clues. The story bible becomes the persistent context used by later planning, memory, anchor, and generation modules.

3.3. DCEE-Tree Planning

DCEE-Tree Planning explores multiple Desire-Conflict routes instead of committing to a single narrative path. Each route is expanded into event-chain candidates, and every event must be causally related to the conflict, visually drawable, and capable of changing the protagonist's emotion. Routes are scored by causal coherence, target-emotion alignment, event richness, visual evidentiality, diversity, and narrative coherence. $\mathcal{R}_j = (d, c_j, \mathcal{E}_{1:K}^{(j)}, e^*)$, where d denotes the central Desire, c_j is the j -th Conflict, and $\mathcal{E}_{1:K}^{(j)}$ is the corresponding Event Chain leading to the target ending emotion e^* .

$$\text{Score}(\mathcal{R}_j) = \alpha s_{\text{causal}} + \beta s_{\text{emotion}} + \gamma s_{\text{richness}} + \delta s_{\text{evidence}} + \eta s_{\text{diversity}} + \zeta s_{\text{coherence}}.$$

Here, the six s terms correspond to the criteria listed above, and $\alpha, \beta, \gamma, \delta, \eta$, and ζ are weighting coefficients. The highest-scoring route becomes the selected DCEE plan.

3.4. Entity-Event Canonicalization

Entity-event canonicalization resolves ambiguous references, normalizes recurring entities, and converts abstract events into drawable visual obligations. For example, an abstract event such as "he loses everything" is rewritten into visible evidence such as an empty-handed protagonist, a broken clue, or an object drifting away. The canonical storyboard records subjects, objects, actions, world context, emotion cues, and evidence objects for every frame.

3.5. Emotion-World Trajectory Planning

merely recent history. The retrieved memory set is defined as

$$M_k = \text{TopR Score}(m_i, q_k)_{m_i \in M_{<k}}$$

where q_k denotes the event-emotion query at frame k , $M_{<k}$ contains the memories stored from previous frames, and TopR selects the r highest-scoring memories. The relevance score is computed as

$$\text{Score}(m_i, q_k) = w_1 s_{\text{event}}(m_i, q_k) + w_2 s_{\text{conflict}}(m_i, q_k) + w_3 s_{\text{cause}}(m_i, q_k) + w_4 s_{\text{evidence}}(m_i, q_k) + w_5 s_{\text{id/world}}(m_i, q_k)$$




Figure 2. Decoder-encoder-decoder structure of the DCEE Butterfly Visual Controller.

The emotion-world trajectory planner assigns each frame an emotion state, intensity, facial cue, body cue, lighting, color palette, weather, camera shot, and composition. Emotions are not treated as labels alone: the planner records emotion_evidence and must_show fields that specify which objects or events must be visible to explain the current affective state.

3.6. Causal Sink Memory and Adaptive Context Selection

Causal sink memory stores the information needed for future story coherence and emotional causality: identity, desire, conflict, event, clue/evidence, world, and emotion-cause memory. At each frame, adaptive context selection retrieves memories using an event-emotion query, so the model conditions on causally relevant history rather than

Here, the five s terms measure event, conflict, emotion-cause, clue/evidence, and identity/world relevance, respectively, and $w_1 \dots, w_5$ are their weighting coefficients.

3.7. DCEE Anchor Bank

The DCEE anchor bank maintains persistent character, object/evidence, and world anchors. Character anchors preserve appearance and signature attributes; evidence anchors preserve causally important objects; and world anchors preserve recurring environments. This extends subject consistency to evidence-object and world-state consistency.

3.8. DCEE Butterfly Visual Controller

The DCEE Butterfly Visual Controller bridges symbolic narrative causality and diffusion-based generation. Directly placing the full storyboard into a long prompt is unreliable

because key evidence may be truncated or entangled with identity, world, and emotion constraints. The controller therefore uses a decoder-encoder-decoder structure: the first decoder expands the current DCEE frame into visual hypotheses, the central encoder fuses them with causal memory and anchor information, and the second decoder outputs factorized Character, World, Emotion, Event, and Evidence controls.

Figure 2 shows how the controller converts a DCEE frame into factorized cross-attention tokens. The Character branch preserves identity, the World branch controls scene and atmosphere, the Emotion branch controls facial expression and posture, the Event branch specifies the action, and the Evidence branch specifies the visible clue or object that explains the emotion.

3.9. Factorized Cross-Attention Control for SDXL

Each control branch is encoded as branch-specific tokens and appended to the original SDXL prompt embedding:

$$H_{\text{afig}} = \left[H_{\text{prompt}}; H_{\text{char}}; H_{\text{world}}; H_{\text{emotion}}; H_{\text{event}}; H_{\text{evidence}} \right].$$

Here, H_{prompt} denotes the original SDXL prompt embedding, while the remaining terms denote the branch-specific token embeddings for Character, World, Emotion, Event, and Evidence; the semicolon indicates token-wise concatenation. This factorization reduces prompt truncation and lets retry generation strengthen a weak branch, such as Event or Evidence, without rewriting the whole prompt.

3.10. Auto-regressive Generation, Reranking, and Retry

Frames are generated autoregressively. For each frame, the system samples multiple candidates and reranks them using image quality, identity consistency, event grounding, evidence visibility, emotion visibility, event-emotion causal consistency, world consistency, and story coherence:

$$S = \lambda_{\text{id}} S_{\text{id}} + \lambda_{\text{q}} S_{\text{quality}} + \lambda_{\text{e}} S_{\text{event}} + \lambda_{\text{v}} S_{\text{evidence}} + \lambda_{\text{m}} S_{\text{emotion}} + \lambda_{\text{c}} S_{\text{causal}} + \lambda_{\text{w}} S_{\text{world}} + \lambda_{\text{s}} S_{\text{story}}$$

Here, the eight S terms correspond to the criteria listed above, and the λ terms are their respective weighting coefficients. The highest-scoring candidate is selected for each frame; If the best candidate falls below event/evidence/emotion thresholds, the weak branch is strengthened and generation is retried.

Table 1. Comparison between prior history-/subject-consistency approaches and DCEE-CausalVerse.

Aspect	Prior history- or subject-consistent methods	DCEE-CausalVerse	Reason for the design choice
Primary goal	Improve temporal coherence or preserve recurring subjects across frames.	Justify a target ending emotion through a causally grounded event chain and visible evidence.	Ending-controllable storytelling requires a believable cause for the final affective state, not only visual consistency.
Input assumption	The caption/story sequence is usually already given.	The event chain is planned from an initial premise, protagonist, and target ending emotion.	Our setting requires the system to decide what should happen before rendering the frames.
History use	Condition on recent or salient text-image history.	Retrieve causal sink memory according to event-emotion relevance.	The current frame uses the history that explains the current emotional state, rather than simply the most recent frame.
Control unit	Caption, subject prompt, or generic history feature.	Factorized Character, World, Emotion, Event, and Evidence controls.	Factorization separates identity preservation from action, setting, emotion, and causal evidence.
Event modeling	Events are often implicit in the caption or absorbed into prompt text.	Events are explicit DCEE units that are planned, canonicalized, grounded, and evaluated per frame.	Explicit event units make it possible to verify whether the generated image depicts the intended action.
Evidence modeling	Evidence objects or clues are usually not modeled separately.	Evidence is a separate branch specifying the visible cause, clue, or object that supports the protagonist emotion.	The target ending emotion is only narratively convincing when its cause is visible.
Generation controller	Prompt generator, history adapter, or identity-conditioning module.	Decoder-encoder-decoder Butterfly controller with causal fusion.	The controller translates symbolic narrative causality into low-level diffusion conditions.
Selection criterion	Image quality, text-image alignment, or subject consistency.	Event grounding, evidence visibility, emotion visibility, event-emotion causal consistency, identity consistency, and world consistency.	Candidate selection directly reflects the causal storytelling objective.

Method	Frame 1	Frame 2	Frame 3	Frame 4	Frame 5	Frame 6
Ours (DCEE-CausalVerse)						
SDXL Prompt-only						
IP-Adapter						
StoryDiffusion (emotion)						
StoryDiffusion (no emotion)						

Figure 3. Qualitative comparison using the shared six-frame storyboard.

Generated story abstract :

"In a deep forest, a white bear anxiously searches for his lost honey jar. He pushes through dense underbrush and tangled foliage, growing frustrated as the search continues. When he reaches a serene lake, hope returns as he spots the honey jar at the water's edge. He approaches the lake, retrieves the honey jar, and finally enjoys sweet honey beside the tranquil water with relief, joy, and contentment."

Generated full story :

1. The white bear searches through dense underbrush, feeling anxious about his lost honey jar. 2. The white bear pushes through tangled foliage, growing increasingly frustrated in his search for the honey jar. 3. The white bear emerges from the tangled foliage, spotting the serene lake ahead with a mix of hope and anxiety. 4. The white bear approaches the lake's edge, feeling a surge of relief upon seeing his honey jar. 5. The white bear happily savors honey beside the serene lake, feeling pure joy at last. 6. The white bear smiles contentedly beside the serene lake, fully enjoying the sweet honey.

4. Experiments

3.11. Method Comparison and Algorithmic Summary

DCEE-CausalVerse differs from prior history- and subject-consistency methods in objective and control granularity. Prior methods mainly render a given sequence consistently, whereas our framework treats the ending emotion as a causal target that must be explained by visible events and evidence. .

4.1. Datasets and Input Construction

We evaluate DCEE-CausalVerse on six-frame visual storytelling examples and a counterfactual ending-control setting in which the same premise can be paired with different target ending emotions. Each input contains a premise, protagonist profile, target ending emotion, and optional reference description; the system outputs a DCEE plan, storyboard, visual control packets, final image sequence, and evaluation logs.

4.2. Baselines and Metrics

We compare DCEE-CausalVerse with Direct LLM + SDXL, SDXL Prompt-only, LLM Storyboard + SDXL, ViSTA-style history conditioning, ConsiStory/StoryDiffusion-style subject-consistent generation, and ablations without the Evidence branch or Tree Planning. For controlled image-level comparison, baselines without a causal planner are conditioned on the same six-frame storyboard. We report CLIP-T, CLIP-I (Radford et al. 2021), and TIFA in the main paper; Appendix C lists the full baseline and metric definitions.

4.3. Implementation Details

All experiments are conducted on a Quad NVIDIA A100 GPU. We use Qwen2-VL-72B (Wang et al. 2024) as the MLLM for rationale and caption generation, with inference accelerated by the vLLM framework (Kwon et al. 2023). In addition, through model lightweighting, it is possible to implement the model without high-end GPUs such as the general computer GPUs RTX 3050 and 3090ti. The implementation uses an LLM for DCEE planning and canonicalization, SDXL as the base image generator, and lightweight PyTorch modules for branch-specific cross-attention control tokens. The pipeline writes the selected DCEE plan, storyboard, visual control packets, memory log, evaluation log, final story, and contact sheet, making planning-to-rendering errors auditable.

4.4. Quantitative Results

We report CLIP-T, CLIP-I, and TIFA. CLIP-T measures frame-caption alignment, CLIP-I measures cross-frame visual consistency, and TIFA measures text-image faithfulness through question answering. We do not report FID (Heusel et al. 2017) because this setting evaluates a modular inference pipeline on story-level examples rather than a trained image distribution.

All baselines are conditioned on the same six-frame storyboard when they do not perform independent causal planning. Table 2 shows that DCEE-CausalVerse obtains the highest CLIP-I score (0.6904), indicating stronger cross-frame visual consistency. It does not yet outperform baselines on CLIP-T or TIFA, showing that local object/action faithfulness and explicit evidence rendering remain important areas for improvement. Figure 3 shows the qualitative comparison.

4.5. Human Evaluation

We follow the compact reporting style of ViSTA and summarize human evaluation using Textual Alignment and Image Alignment. Ten participants compared five anonymized sequences, where Method A denotes DCEE-CausalVerse. Textual Alignment averages Q3-Q6 (caption reflection, object/action visibility, ending-emotion delivery, and emotion-cause explainability). Image Alignment averages Q2 and Q8 (reference adherence and overall visual quality). Table 3 shows that Method A achieves the highest Textual Alignment (4.60) and competitive Image Alignment (3.85). Detailed questionnaires, Likert scores, and preference counts are in Appendix A.

Table 2. Automatic quantitative results averaged over the selected six frames.

Method	CLIP-T $\uparrow$	CLIP-I $\uparrow$	TIFA $\uparrow$
Ours (DCEE-CausalVerse)	0.3352	0.6904	0.3667
SDXL Prompt-only	0.3456	0.6274	0.6667
IP-Adapter	0.3422	0.6699	0.6667
StoryDiffusion (emotion)	0.3418	0.6699	0.5000
StoryDiffusion (no emotion)	0.3479	0.6673	0.6333

Table 3. ViSTA-style human evaluation summary. Scores are averaged over ten participants. Method A denotes DCEE-CausalVerse.

Model	Textual Align. $\uparrow$	Image Align. $\uparrow$
Method A (Ours)	4.60	3.85
Method B	2.88	3.25
Method C	3.80	3.90
Method D	2.33	2.55
Method E	2.05	2.55

5. Discussion

DCEE-CausalVerse is designed to make visual storytelling more causally explainable. The automatic results show stronger cross-frame consistency, and the human study shows that viewers prefer our method on story-grounded textual alignment. At the same time, the lower CLIP-T/TIFA scores indicate that local evidence rendering remains challenging, especially for small objects and fine-grained actions. This suggests that future work should combine causal planning with stronger local grounding modules or trainable evidence-aware diffusion adapters. A full component-wise ablation of DCEE-Tree Planning, causal sink memory, anchor bank, and the Evidence branch is left to future work rather than reported as an incomplete ablation plan.

6. Conclusion

We presented DCEE-CausalVerse, an event-grounded framework for ending-controllable visual storytelling. The method decomposes a target ending emotion into Desire, Conflict, Event Chain, and Ending Emotion, explores multiple causal routes, canonicalizes events into drawable evidence, maintains causal memory and anchors, and uses a decoder-encoder-decoder Butterfly controller to produce factorized diffusion controls. Results indicate stronger cross-frame consistency and higher human preference for story-grounded textual alignment, while also revealing the remaining challenge of local event/evidence faithfulness. Overall, our work moves visual storytelling from sequence consistency alone toward causal, interpretable, and emotionally controllable visual narratives.

References

- Ahn, D.; Kim, D.; Song, G.; Kim, S. H.; Lee, H.; Kang, D.; and Choi, J. 2023. Story Visualization by Online Text Augmentation with Context Memory. In Proceedings of the IEEE/CVF International Conference on Computer Vision.
- Dong, S.; Shaheen, I.; Shen, M.; Mallick, R.; and Bargal, S. A. 2026. ViSTA: Visual Storytelling Using Multi-Modal Adapters for Text-to-Image Diffusion Models. In Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision.
- Fan, A.; Lewis, M.; and Dauphin, Y. 2018. Hierarchical Neural Story Generation. In Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers).
- Feng, Z.; Ren, Y.; Yu, X.; Feng, X.; Tang, D.; Shi, S.; and Qin, B. 2023. Improved Visual Story Generation with Adaptive Context Modeling. In Findings of the Association for Computational Linguistics.
- Ghosh, D.; Hajishirzi, H.; and Schmidt, L. 2023. GenEval: An Object-Focused Framework for Evaluating Text-to-Image Alignment. In Advances in Neural Information Processing Systems.
- He, H.; Yang, H.; Tuo, Z.; Zhou, Y.; Wang, Q.; Zhang, Y.; Liu, Z.; Huang, W.; Chao, H.; and Yin, J. 2025. DreamStory: Open-Domain Story Visualization by LLM-Guided Multi-Subject Consistent Diffusion. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 47(12): 11874–11891.
- Heusel, M.; Ramsauer, H.; Unterthiner, T.; Nessler, B.; and Hochreiter, S. 2017. GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium. In Advances in Neural Information Processing Systems.
- Ho, J.; Jain, A.; and Abbeel, P. 2020. Denoising Diffusion Probabilistic Models. In Advances in Neural Information Processing Systems.
- Hu, Y.; Liu, B.; Kasai, J.; Wang, Y.; Ostendorf, M.; Krishna, R.; and Smith, N. A. 2023. TIFA: Accurate and Interpretable Text-to-Image Faithfulness Evaluation with Question Answering. In Proceedings of the IEEE/CVF International Conference on Computer Vision.
- Huang, T.-H. K.; Ferraro, F.; Mostafazadeh, N.; Misra, I.; Agrawal, A.; Devlin, J.; Girshick, R.; He, X.; Kohli, P.; Batra, D.; Zitnick, C. L.; Parikh, D.; Vanderwende, L.; Galley, M.; and Mitchell, M. 2016. Visual Storytelling. In Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies.
- Kwon, W.; Li, Z.; Zhuang, S.; Sheng, Y.; Zheng, L.; Yu, C. H.; Gonzalez, J. E.; Zhang, H.; and Stoica, I. 2023. Efficient Memory Management for Large Language Model Serving with PagedAttention. In Proceedings of the 29th Symposium on Operating Systems Principles.
- Li, B.; and Lukasiewicz, T. 2022. Learning to Model Multimodal Semantic Alignment for Story Visualization. In Findings of the Association for Computational Linguistics: EMNLP 2022.
- Li, Y.; Gan, Z.; Shen, Y.; Liu, J.; Cheng, Y.; Wu, Y.; Carin, L.; Carlson, D.; and Gao, J. 2019. StoryGAN: A Sequential Conditional GAN for Story Visualization. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition.
- Lin, Z.; Pathak, D.; Li, B.; Li, J.; Xia, X.; Neubig, G.; Zhang, P.; and Ramanan, D. 2024. Evaluating Text-to-Visual Generation with Image-to-Text Generation. In

- Proceedings of the European Conference on Computer Vision.
- Liu, C.; Wu, H.; Zhong, Y.; Zhang, X.; Wang, Y.; and Xie, W. 2024. Intelligent Grimm: Open-Ended Visual Storytelling via Latent Diffusion Models. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition.
- Liu, T.; Wang, K.; Li, S.; van de Weijer, J.; Khan, F.; Yang, S.; Wang, Y.; Yang, J.; and Cheng, M.-M. 2025. One-Prompt-One-Story: Free-Lunch Consistent Text-to-Image Generation Using a Single Prompt. In Proceedings of the International Conference on Learning Representations.
- Maharana, A.; Hannan, D.; and Bansal, M. 2022. StoryDALL-E: Adapting Pretrained Text-to-Image Transformers for Story Continuation. In Proceedings of the European Conference on Computer Vision.
- Pan, X.; Qin, P.; Li, Y.; Xue, H.; and Chen, W. 2024. Synthesizing Coherent Story With Auto-Regressive Latent Diffusion Models. In Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision.
- Podell, D.; English, Z.; Lacey, K.; Blattmann, A.; Dockhorn, T.; Müller, J.; Penna, J.; and Rombach, R. 2024. SDXL: Improving Latent Diffusion Models for High-Resolution Image Synthesis. In Proceedings of the International Conference on Learning Representations.
- Radford, A.; Kim, J. W.; Hallacy, C.; Ramesh, A.; Goh, G.; Agarwal, S.; Sastry, G.; Askell, A.; Mishkin, P.; Clark, J.; Krueger, G.; and Sutskever, I. 2021. Learning Transferable Visual Models From Natural Language Supervision. arXiv preprint arXiv:2103.00020.
- Rahman, T.; Lee, H.-Y.; Ren, J.; Tulyakov, S.; Mahajan, S.; and Sigal, L. 2023. Make-A-Story: Visual Memory Conditioned Consistent Story Generation. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition.
- Rashkin, H.; Celikyilmaz, A.; Choi, Y.; and Gao, J. 2020. PlotMachines: Outline-Conditioned Generation with Dynamic Plot State Tracking. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing.
- Rombach, R.; Blattmann, A.; Lorenz, D.; Esser, P.; and Ommer, B. 2022. High-Resolution Image Synthesis With Latent Diffusion Models. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition.
- Shen, X.; and Elhoseiny, M. 2025. StoryGPT-V: Large Language Models as Consistent Story Visualizers. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition.
- Song, T.; Cao, J.; Wang, K.; Liu, B.; and Zhang, X. 2024. Causal-Story: Local Causal Attention Utilizing Parameter-Efficient Tuning for Visual Story Synthesis. In Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing.
- Tambwekar, P.; Dhuliawala, M.; Martin, L. J.; Mehta, A.; Harrison, B.; and Riedl, M. O. 2019. Controllable Neural Story Plot Generation via Reward Shaping. In Proceedings of the International Joint Conference on Artificial Intelligence.
- Tewel, Y.; Kaduri, O.; Gal, R.; Kasten, Y.; Wolf, L.; Chechik, G.; and Atzmon, Y. 2024. Training-Free Consistent Text-to-Image Generation. ACM Transactions on Graphics 43(4): 52:1–52:18.
- Wang, P.; Bai, S.; Tan, S.; Wang, S.; Fan, Z.; Bai, J.; Chen, K.; Liu, X.; Wang, J.; Ge, W.; Fan, Y.; Dang, K.; Du, M.; Ren, X.; Men, R.; Liu, D.; Zhou, C.; Zhou, J.; and Lin, J. 2024. Qwen2-VL: Enhancing Vision-Language Model's Perception of the World at Any Resolution. arXiv preprint arXiv:2409.12191.
- Yang, S.; Ge, Y.; Li, Y.; Chen, Y.; Ge, Y.; Shan, Y.; and Chen, Y.-C. 2025. SEED-Story: Multimodal Long Story Generation with Large Language Model. In Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops.
- Yao, L.; Peng, N.; Weischedel, R.; Knight, K.; Zhao, D.; and Yan, R. 2019. Plan-and-Write: Towards Better Automatic Storytelling. In Proceedings of the AAAI Conference on Artificial Intelligence.
- Ye, H.; Zhang, J.; Liu, S.; Han, X.; and Yang, W. 2023. IP-Adapter: Text Compatible Image Prompt Adapter for Text-to-Image Diffusion Models. arXiv preprint arXiv:2308.06721.
- Zhou, Y.; Zhou, D.; Cheng, M.-M.; Feng, J.; and Hou, Q. 2024. StoryDiffusion: Consistent Self-Attention for Long-Range Image and Video Generation. In Advances in Neural Information Processing Systems.